

Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning:
pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
import pandas.util.testing as tm
```

In [2]:

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

In [3]:

```
df = pd.read_csv('/content/drive/My Drive/iris_model/Iris.csv', index_col=False)
```

In [4]:

```
df.drop('Id', axis=1, inplace=True)
```

In [5]:

```
df.head()
```

Out[5]:

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SepalLengthCm    150 non-null   float64
1   SepalWidthCm     150 non-null   float64
2   PetalLengthCm    150 non-null   float64
3   PetalWidthCm     150 non-null   float64
4   Species          150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

In [7]:

```
df.describe().T
```

Out[7]:

	count	mean	std	min	25%	50%	75%	max
SepalLengthCm	150.0	5.843333	0.828066	4.3	5.1	5.80	6.4	7.9
SepalWidthCm	150.0	3.054000	0.433594	2.0	2.8	3.00	3.3	4.4
PetalLengthCm	150.0	3.758667	1.764420	1.0	1.6	4.35	5.1	6.9
PetalWidthCm	150.0	1.198667	0.763161	0.1	0.3	1.30	1.8	2.5

In [8]:

```
df.isnull().sum()
```

Out[8]:

```
SepalLengthCm    0
SepalWidthCm     0
PetalLengthCm    0
PetalWidthCm     0
Species          0
dtype: int64
```


In [9]:

```
setosa = df[df.Species == 'Iris-setosa']
versicolor = df[df.Species == 'Iris-versicolor']
virginica = df[df.Species == 'Iris-virginica']
```

Visual EDA

In [10]:

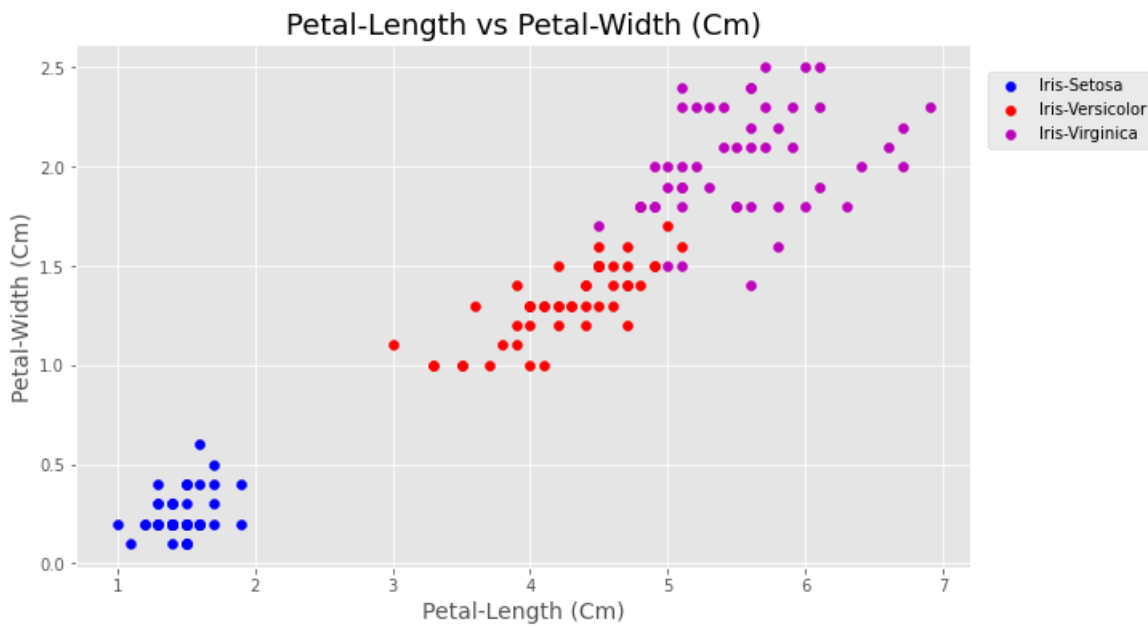
```
plt.figure(figsize=(10,6))
plt.style.use('ggplot')
plt.scatter(setosa['SepalLengthCm'], setosa['SepalWidthCm'], c = 'b', label = 'Iris-Setosa')
plt.scatter(versicolor['SepalLengthCm'], versicolor['SepalWidthCm'], c = 'r', label = 'Iris-Versicolor')
plt.scatter(virginica['SepalLengthCm'], virginica['SepalWidthCm'], c = 'm', label = 'Iris-Virginica')
plt.xlabel('Sepal-Length (Cm)', fontsize = 14)
plt.ylabel('Sepal-Width (Cm)', fontsize = 14)
plt.title('Sepal-Length vs Sepal-Width (Cm)', fontsize = 18)
plt.legend(loc = (1.02, 0.8))
plt.show()
```



Sepal-Length (Cm)


In [11]:

```
plt.figure(figsize=(10,6))
plt.style.use('ggplot')
plt.scatter(setosa['PetalLengthCm'], setosa['PetalWidthCm'], c = 'b', label = 'Iris-Setosa')
plt.scatter(versicolor['PetalLengthCm'], versicolor['PetalWidthCm'], c = 'r', label = 'Iris-Versicolor')
plt.scatter(virginica['PetalLengthCm'], virginica['PetalWidthCm'], c = 'm', label = 'Iris-Virginica')
plt.xlabel('Petal-Length (Cm)', fontsize = 14)
plt.ylabel('Petal-Width (Cm)', fontsize = 14)
plt.title('Petal-Length vs Petal-Width (Cm)', fontsize = 18)
plt.legend(loc = (1.02, 0.8))
plt.show()
```



In [12]:

```
df.hist(edgecolor = 'black', linewidth = 1.2)
plt.gcf().set_size_inches(10, 6)
plt.show()
```



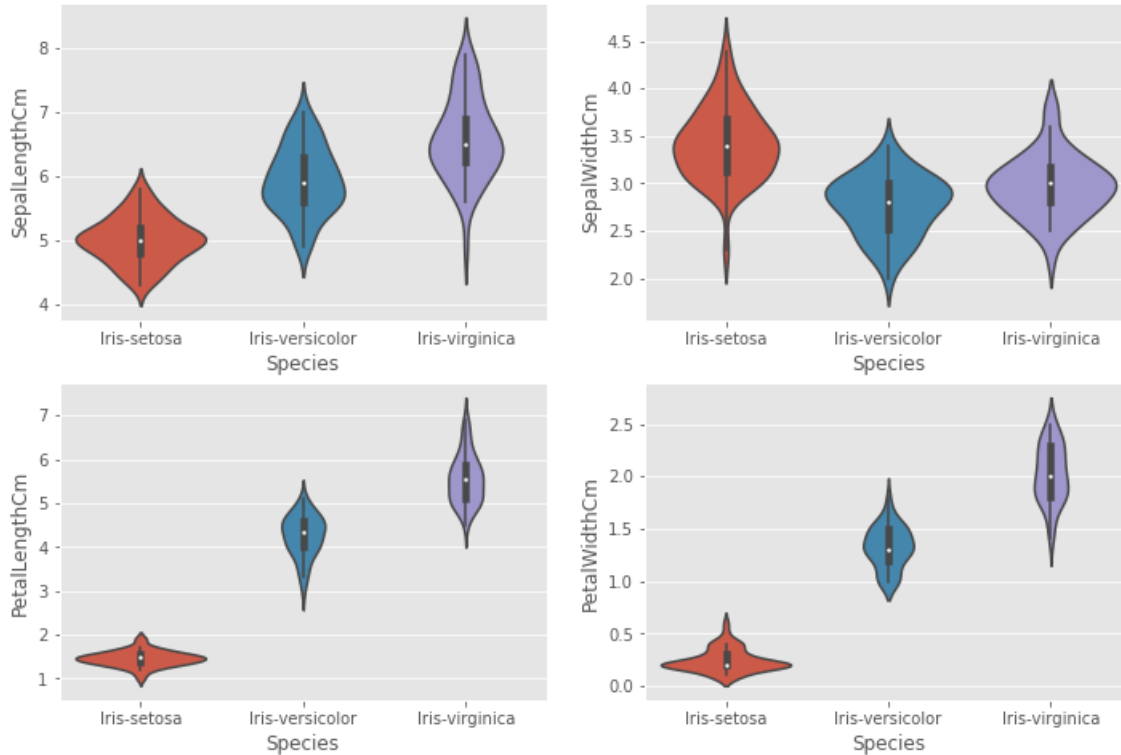
In [13]:

```
plt.figure(figsize=(12, 8))
```

```

plt.figure(figsize=(12,8))
plt.subplot(2,2,1)
sns.violinplot(x='Species', y='SepalLengthCm', data=df)
plt.subplot(2,2,2)
sns.violinplot(x='Species', y='SepalWidthCm', data=df)
plt.subplot(2,2,3)
sns.violinplot(x='Species', y='PetalLengthCm', data=df)
plt.subplot(2,2,4)
sns.violinplot(x='Species', y='PetalWidthCm', data=df)
plt.show()

```




In [14]:

```

sns.boxplot(data=df)
plt.show()

```



In [15]:

```
df.shape
```

Out[15]:

(150, 5)

Building SkLearn Model

In [16]:

```
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.metrics import mean_squared_error, accuracy_score, classification_report,
confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
```

Train Test Split


In [17]:

```
X = df.drop('Species', axis = 1)
y = df['Species']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 42, stratify = y)
```

Feature Selection

In [18]:


```
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(), annot = True, cmap = 'RdYlGn', square=True)
plt.yticks(rotation = 0)
plt.xticks(rotation = 90)
plt.show()
```



In [19]:


```
model = RandomForestClassifier(n_estimators=100, random_state=0)
model.fit(X, y)
imp_features = pd.Series(model.feature_importances_, index = X.columns).sort_values(ascending=False)
# print(imp_features)
plt.plot(imp_features, color = 'm')
plt.show()
```





In [20]:

```
from sklearn.preprocessing import LabelEncoder
g = y.copy()
le = LabelEncoder().fit(g)
# print(le.classes_)
encoded_y = le.transform(g)
from sklearn.linear_model import LassoCV
names = X.columns
lasso = LassoCV(cv=3)
features = lasso.fit(X, encoded_y).coef_
plt.plot(features, color = 'blue')
plt.xticks(range(len(names)), names, rotation = 0)
plt.figure(figsize=(7,4))
plt.show()
```



<Figure size 504x288 with 0 Axes>

Checking Accuracy of Each Model

In [21]:

```
accuracy = []
error = []
classifiers = ['Linear_SVC', 'Radial_SVC', 'KNN', 'Logistic_Regression', 'Decision_Tree', 'Random_Forest']
models = [SVC(gamma='scale', kernel='linear'), SVC(gamma='auto', kernel='rbf'),
           KNeighborsClassifier(n_neighbors=5),
           LogisticRegression(solver='liblinear', multi_class='auto'), DecisionTreeClassifier(), RandomForestClassifier(n_estimators=100)]
for i in models:
    model = i
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy.append(accuracy_score(y_pred, y_test))
    y_pred_encoded = LabelEncoder().fit_transform(y_pred)
    y_test_encoded = LabelEncoder().fit_transform(y_test)
    error.append(np.sqrt(mean_squared_error(y_test_encoded, y_pred_encoded)))
d = {'Accuracy' : accuracy, 'RMSE' : error}
score = pd.DataFrame(d, index = classifiers)
score
```

Out[21]:

	Accuracy	RMSE
Linear_SVC	1.000000	0.000000
Radial_SVC	0.977778	0.149071
KNN	0.977778	0.149071
Logistic_Regression	0.911111	0.298142
Decision_Tree	0.977778	0.149071
Random_Forest	0.888889	0.333333

Hyper-Parameter Tuning (KNN)

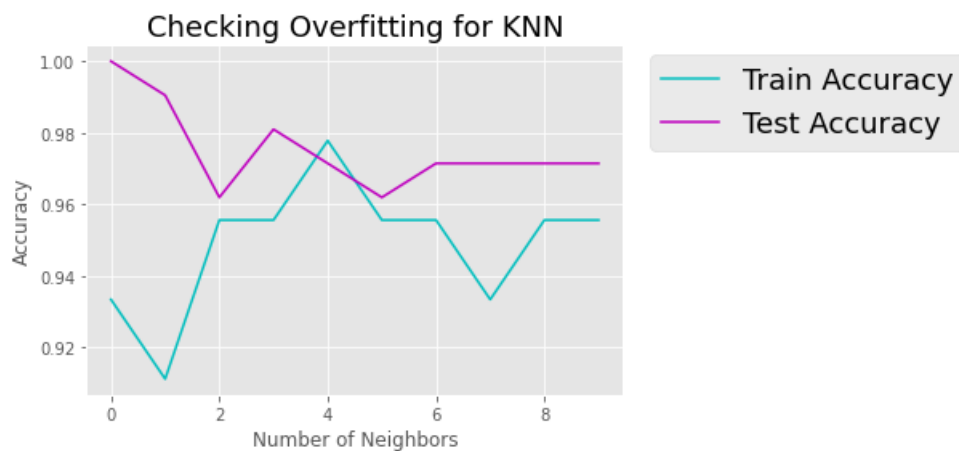
In [22]:

```
param_grid = {'n_neighbors' : np.arange(1, 51)}
knn_model = KNeighborsClassifier()
knn_cv = GridSearchCV(knn_model, param_grid, cv = 5, scoring='accuracy')
knn_cv.fit(X, y)
print(knn_cv.best_params_)
print(knn_cv.best_score_)
```

```
{'n_neighbors': 6}
0.9800000000000001
```

In [23]:

```
neighbors = np.arange(1, 11)
test_accuracy = np.empty(len(neighbors))
train_accuracy = np.empty(len(neighbors))
for i, k in enumerate(neighbors):
    model = KNeighborsClassifier(n_neighbors=k)
    model.fit(X_train, y_train)
    test_accuracy[i] = model.score(X_test, y_test)
    train_accuracy[i] = model.score(X_train, y_train)
plt.plot(train_accuracy, label='Train Accuracy', color = 'c')
plt.plot(test_accuracy, label='Test Accuracy', color = 'm')
plt.title('Checking Overfitting for KNN', fontsize = 18)
plt.xlabel('Number of Neighbors')
plt.ylabel('Accuracy')
plt.legend(fontsize=18, loc = (1.05, 0.7))
plt.figure(figsize=(12,6))
plt.show()
```



<Figure size 864x432 with 0 Axes>

Choosing Linear_SVC

In [24]:

```
Linear_SVC (C=1.0, gamma=0.5, kernel='linear', ...)
```

```

model = SVC(gamma='auto', kernel='linear')
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print(f'Root Mean Square Error: {mean_squared_error(y_test_encoded, y_pred_encoded)}')
print(f'Model Accuracy: {round(accuracy_score(y_test, y_pred), 3) * 100}%')
print('=====')
print(f'Here is the "Classification Report"')
print(classification_report(y_test, y_pred))
print(f'\nHere is the "Confusing Matrix"\n')
print(confusion_matrix(y_test, y_pred))
print('=====')

```

Root Mean Square Error: 0.1111111111111111

Model Accuracy: 100.0%

=====

Here is the "Classification Report"

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	15
Iris-versicolor	1.00	1.00	1.00	15
Iris-virginica	1.00	1.00	1.00	15
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

Here is the "Confusing Matrix"

```

[[15  0  0]
 [ 0 15  0]
 [ 0  0 15]]

```

=====

Creating pipeline

In [25]:

```

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
steps = [('model', SVC(kernel='linear', gamma = 'auto'))]
pipeline = Pipeline(steps)
pipeline.fit(X_train, y_train)
y_pred = pipeline.predict(X_test)
print('{}%'.format(accuracy_score(y_test, y_pred)*100))

```

100.0%

Saving and Loading the Model

In [26]:

```

import joblib
#save the model
joblib.dump(pipeline, 'iris_model.pkl')
new_model = joblib.load('iris_model.pkl')
new_model.predict(X_test)
print('{}%'.format(accuracy_score(y_test, y_pred)*100))

```

100.0%

Building TensorFlow/Keras Model

In [27]:

```

from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import Dense

```



```

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import to_categorical

df = pd.read_csv('/content/drive/My Drive/iris_model/Iris.csv', index_col=False)
X = df.drop(['Species', 'Id'], axis = 1).values
y = df['Species'].values
y = LabelEncoder().fit_transform(y)
y = to_categorical(y)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)

```

In [28]:

```

# Build the model
model = Sequential()

model.add(Dense(25, input_shape=(4,), activation='relu'))
model.add(Dense(25, activation='relu'))
model.add(Dense(3, activation='softmax'))

# Adam optimizer with learning rate of 0.001
optimizer = Adam(lr=0.001)
model.compile(optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

print('Neural Network Model Summary: ')
print(model.summary())

```

Neural Network Model Summary:

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 25)	125
dense_1 (Dense)	(None, 25)	650
dense_2 (Dense)	(None, 3)	78
Total params: 853		
Trainable params: 853		
Non-trainable params: 0		

None

In [29]:

```
model.fit(X_train, y_train, verbose=0, batch_size=5, epochs=200)
```

Out[29]:

<tensorflow.python.keras.callbacks.History at 0x7f44d198f668>

In [30]:

```

results = model.evaluate(X_test, y_test)

print('Test Loss: {:.2f}'.format(results[0]))
print('Test Accuracy: {:.2f}'.format(results[1]))

```

1/1 [=====] - 0s 2ms/step - loss: 0.0115 - accuracy: 1.0000
Test Loss: 0.01
Test Accuracy: 1.00

In [31]:

```
model.save('keras_model.h5')
```

In [32]:

```
from tensorflow.keras.models import load_model
```

In [33]:

```
new_model = load_model('keras_model.h5')
```

In [34]:

```
res = new_model.evaluate(X_test, y_test)
print('Test Loss: {:.2f}'.format(res[0]))
print('Test Accuracy: {:.2f}'.format(res[1]))
```

1/1 [=====] - 0s 1ms/step - loss: 0.0115 - accuracy: 1.0000

Test Loss: 0.01

Test Accuracy: 1.00

In [34]: